

DATA STRUCTURES (AND ALGORITHMS)

Heap. ADT Priority Queue.

Lect. PhD. Diana-Lucia Miholca

2023 - 2024



Babeş - Bolyai University
Faculty of Mathematics and Computer Science

In the previous lecture...

- Binary Search Trees

Today

- Binary Heap
- ADT Priority Queue

Binary Heap - Introduction



A binary heap is a data structure whose elements are stored in a dynamic array, but that can be visualized as a binary tree.



A binary heap is a data structure that is particularly efficient for representing Priority Queues.


Heap - Example



Assume that we have the following array:

1	2	3	4	5	6	7	8	9	10	11	12	13	14
3	6	14	11	19	28	21	34	22	23	67	85	44	31

- We can visualize this array as an **almost complete** binary tree whose root is the first element, its children are the next two elements and so on.



Heap - Properties



If the elements of the array are: $a_1, a_2, a_3, \dots, a_n$, we know that:

- a_1 is the root of the heap
- for the element at index i , its children are on indexes $2 * i$ and $2 * i + 1$ (if $\leq n$)
- for the element at index i ($i > 1$), its parent is at index $\lceil i/2 \rceil$

Heap - Definition



A **binary heap** is an array that can be visualized as a binary tree having a **heap structure** and a **heap property**.



Heap structure: the binary tree is *almost complete* (all the levels being completely filled, excepting the last one, which is filled from left to right)



Heap property: $a_i \geq a_{2*i}$ (if $2 * i \leq n$) and $a_i \geq a_{2*i+1}$ (if $2 * i + 1 \leq n$)



The \geq relation between a node and its children can be replaced by \leq . By doing so, we will have a *min-heap* instead of a *max-heap*.

Heap - Examples



Are the following binary trees heaps?

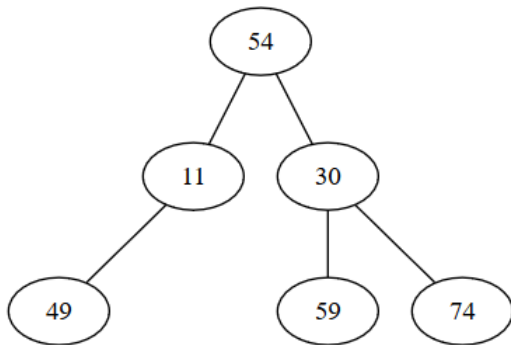
- If yes, specify the relation between a node and its children.
- If not, specify if the problem is with the structure, the property or both.

Heap - Examples



Are the following binary trees heaps?


- If yes, specify the relation between a node and its children.
- If not, specify if the problem is with the structure, the property or both.



Heap - Examples




Is this a heap?



Heap - Examples



Is this a heap?

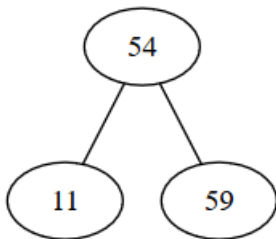


Correct answer: No

Heap - Examples



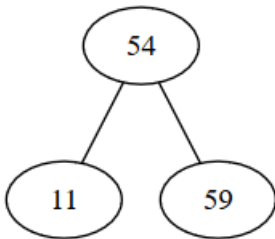
Is this a heap?



Heap - Examples



Is this a heap?




Correct answer: No

Heap - Examples




Is this a heap?



Heap - Examples



Is this a heap?

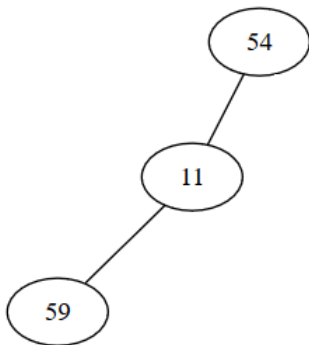


Correct answer: Yes

Heap - Examples




Is this a heap?



Heap - Examples



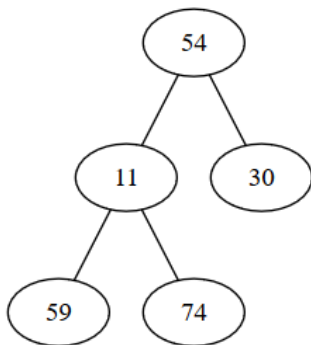
Is this a heap?



Heap - Examples



Is this a heap?



Heap - Examples



Is the following array a valid heap? If not, transform it into a valid heap by swapping two elements.

[70, 10, 50, 7, 1, 33, 3, 8]

Heap - Height



What is the height of a heap with n elements?



What is the height of a heap with n elements?



The height of a heap with n elements is $\lceil \log_2 n \rceil$


Heap - Operations

- Specific operations:
 - **add** a new element to the heap (while keeping both the heap structure and the heap property);
 - **remove** **the root** of the heap (no other element can be removed).

Heap - Add - Example




Consider the following (MAX) heap:




- Let's add the number 55 to the heap.

Heap - Add - Example

- In order to keep the *heap structure*, we will add the new node as the right child of the node 14 (and as the last element of the array in which the elements are kept).



Heap - Add - Example




Heap property is not kept: 14 has as right child 55 and $14 < 55$.



We will perform a *bubble-up* process: we will keep swapping the value of the new node with the value of its parent node, until the heap property is met.

Heap - Add - Example

- When *bubble-up* ends:



Heap - Representation



For the implementation, we assume that we have a MAX-HEAP.



Heap representation:

Heap:

Heap - Representation



For the implementation, we assume that we have a MAX-HEAP.



Heap representation:

Heap:

capacity: Integer

length: Integer

elems: TComp[]



Adding a new element to the heap:

subalgorithm add(heap, e) **is:**

//heap - a heap

//e - the element to be added



Adding a new element to the heap:

subalgorithm add(heap, e) **is:**

//heap - a heap

//e - the element to be added

if heap.length = heap.capacity **then**

 @ resize

end-if

heap.length \leftarrow heap.length + 1

heap.elms[heap.length] \leftarrow e

bubble-up(heap, heap.length)

end-subalgorithm



Bubble-up for restoring the heap property:

subalgorithm bubble-up (heap, p) **is:**

//heap - a heap

//p - position from which we bubble the new node up



Bubble-up for restoring the heap property:

subalgorithm bubble-up (heap, p) **is:**

//heap - a heap

//p - position from which we bubble the new node up

$\text{pos} \leftarrow p$

$\text{elem} \leftarrow \text{heap.elems}[p]$

$\text{parent} \leftarrow [p / 2]$

while $\text{pos} > 1$ **and** $\text{elem} > \text{heap.elems}[\text{parent}]$ **execute**

//move parent down

$\text{heap.elems}[\text{pos}] \leftarrow \text{heap.elems}[\text{parent}]$

$\text{pos} \leftarrow \text{parent}$

$\text{parent} \leftarrow [\text{pos}/2]$

end-while

$\text{heap.elems}[\text{pos}] \leftarrow \text{elem}$

end-subalgorithm



What is the time complexity of *bubble-up*?

Heap - Add



What is the time complexity of *bubble-up*?



$O(\log_2 n)$

Heap - Add



What is the time complexity of *bubble-up*?



$O(\log_2 n)$



What is a best case scenario?

Heap - Add



What is the time complexity of *bubble-up*?



$O(\log_2 n)$



What is a best case scenario?



When the last (newly added) element is already less than or equal to its parent.

Heap - Add



What is the time complexity of *bubble-up*?



$O(\log_2 n)$



What is a best case scenario?



When the last (newly added) element is already less than or equal to its parent.



What is the time complexity of *add*?

Heap - Add



What is the time complexity of *bubble-up*?



$O(\log_2 n)$



What is a best case scenario?



When the last (newly added) element is already less than or equal to its parent.



What is the time complexity of *add*?




$O(\log_2 n)$ *amortized*

Heap - Remove - Example




From a heap we can only remove the root element.




Heap - Remove - Example

- To keep the *heap structure*, when we remove the root, we replace it with the last element on the last level.



Heap - Remove - Example




The *heap property* is violated (the root is no longer the maximum element).



To restore the heap property, we will do a *bubble-down*: the new root element will be swapped with its maximum child, until it becomes a leaf, or until it will be greater than its both children.

Heap - Remove - Example

- When *bubble-down* ends:



Heap - Remove



Deleting from a heap:

function remove(heap) **is:**

//heap - is a heap

//returns the deleted element

Heap - Remove



Deleting from a heap:

function remove(heap) **is:**

//heap - is a heap

//returns the deleted element

if heap.length = 0 **then**

 @ error - empty heap

end-if

deletedElem \leftarrow heap.elems[1]

heap.elems[1] \leftarrow heap.elems[heap.length]

heap.length \leftarrow heap.length - 1

bubble-down(heap, 1)

remove \leftarrow deletedElem

end-function



Bubble-down for restoring the heap property:

subalgorithm bubble-down(heap, p) **is:**

Heap - Remove



Bubble-down for restoring the heap property:

subalgorithm bubble-down(heap, p) **is**:

pos \leftarrow p

elem \leftarrow heap.elems[p]

while pos \leq [heap.length/2] **execute**

maxChild \leftarrow pos*2

if pos*2+1 \leq heap.length **and** heap.elems[2*pos+1] > heap.elems[2*pos] **then**

//it has two children and the right is greater

maxChild \leftarrow pos*2 + 1

end-if

if heap.elems[maxChild] > elem **then**

tmp \leftarrow heap.elems[pos]

heap.elems[pos] \leftarrow heap.elems[maxChild]

heap.elems[maxChild] \leftarrow tmp

pos \leftarrow maxChild

else

pos \leftarrow heap.length + 1 *//to stop the while loop*

end-if

end-while

end-subalgorithm



What is the time complexity of *bubble-down*?

Heap - Remove



What is the time complexity of *bubble-down*?



$O(\log_2 n)$

Heap - Remove



What is the time complexity of *bubble-down*?



$O(\log_2 n)$



What is the time complexity of *remove*?

Heap - Remove



What is the time complexity of *bubble-down*?



$O(\log_2 n)$



What is the time complexity of *remove*?



$O(\log_2 n)$



Applications of heaps:



Graph algorithms

- Used in Dijkstra's algorithm (shortest path) or Prim's algorithm (minimum spanning tree)



Data compression

- Huffman coding



Sorting (Heap-Sort)

- Used by embedded systems such as Linux Kernel



Order statistics

- To find the k th smallest or largest element in an array



Resources allocation

- To efficiently allocate resources in a system, such as memory blocks or CPU time, by processing requests in order of priority



Source: VectorStock.com/21245929



Consider the following queue in front of the Emergency Room.



Who should be the next person checked by the doctor?



The ADT Priority Queue is a container in which each element has an associated *priority*.



In a Priority Queue, the access to the elements is restricted: we can access only the element with the highest priority.



Due to the restricted access, a Priority Queue is said to have a **HPF - Highest Priority First** policy.



In order to work in a more general manner, we can define a relation \mathcal{R} on the set of priorities: $\mathcal{R} : TPriority \times TPriority$.



When we refer to *the element with the highest priority* we mean the highest priority given by the relation \mathcal{R} .

- If the relation $\mathcal{R} = "\geq"$, the element with the *highest priority* is the one for which the value of the priority is the largest.
- If the relation $\mathcal{R} = "\leq"$, the element with the *highest priority* is the one for which the value of the priority is the lowest.

Priority Queue - Domain

- The domain of the ADT Priority Queue:

$\mathcal{PQ} = \{pq \mid pq \text{ is a priority queue with elements } (e, p), e \in T\text{Elem}, p \in T\text{Priority}\}$

- `init (pq, R)`
 - **descr:** creates a new empty priority queue
 - **pre:** R is a relation over the priorities, $R : \text{TPriority} \times \text{TPriority}$
 - **post:** $pq \in \mathcal{PQ}$, pq is an empty priority queue

Priority Queue - Interface

- `destroy(pq)`
 - **descr:** destroys a priority queue
 - **pre:** $pq \in \mathcal{PQ}$
 - **post:** pq has been destroyed

- `push(pq, e, p)`
 - **descr:** pushes (adds) a new element to the priority queue
 - **pre:** $pq \in \mathcal{PQ}, e \in TElem, p \in TPriority$
 - **post:** $pq' \in \mathcal{PQ}, pq' = pq \oplus (e, p)$

Priority Queue - Interface

- **pop(pq)**
 - **descr:** pops (removes) from the priority queue the element with the highest priority. It returns both the element and its priority
 - **pre:** $pq \in \mathcal{PQ}$, pq is not empty
 - **post:** $pop = (e, p)$, $e \in TElem$, $p \in TPriority$, e is the element with the highest priority from pq , p is its priority.
 $pq' \in \mathcal{PQ}$, $pq' = pq \ominus (e, p)$
 - **throws:** an exception if the priority queue is empty.

Priority Queue - Interface

- **top(pq)**
 - **descr:** returns from the priority queue the element with the highest priority and its priority. It does not modify the priority queue.
 - **pre:** $pq \in \mathcal{PQ}$, pq is not empty
 - **post:** $top = (e, p)$, $e \in TElem$, $p \in TPriority$, e is the element with the highest priority from pq , p is its priority.
 - **throws:** an exception if the priority queue is empty.

- `isEmpty(pq)`
 - **descr:** checks if the priority queue is empty (it has no elements)
 - **pre:** $pq \in \mathcal{PQ}$
 - **post:**

$$isEmpty = \begin{cases} true, & \text{if } pq \text{ has no elements} \\ false, & \text{otherwise} \end{cases}$$



Priority queues cannot be iterated, so they don't have an *iterator* operation.

Priority Queue - Representation



Data structures can be used to represent a Priority Queue:

- Dynamic Array
- Linked List
- Binary Heap

Priority Queue - Representation



If we opt for a Dynamic Array or a Linked List we have to decide in which order to store the elements:

- We can keep the elements ordered by their priorities.



Where would you put the element with the highest priority?

Priority Queue - Representation



If we opt for a Dynamic Array or a Linked List we have to decide in which order to store the elements:

- We can keep the elements ordered by their priorities.
 - Where would you put the element with the highest priority?
 - ✓ At the beginning for linked lists and at the end for arrays.
- We also can keep the elements in the order in which they have been inserted.

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push		

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	$\Theta(1)$
pop		

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	$\Theta(1)$
pop	$\Theta(1)$	

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	$\Theta(1)$
pop	$\Theta(1)$	$\Theta(n)$
top		

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	$\Theta(1)$
pop	$\Theta(1)$	$\Theta(n)$
top	$\Theta(1)$	

Priority Queue - Representation



Complexity of the main operations for the two representation options:

Operation	Sorted	Non-sorted
push	$O(n)$	$\Theta(1)$
pop	$\Theta(1)$	$\Theta(n)$
top	$\Theta(1)$	$\Theta(n)$

Priority Queue - Representation on a binary heap



When representing a Priority Queue using a Max Binary Heap:



When an element is pushed to the priority queue

Priority Queue - Representation on a binary heap



When representing a Priority Queue using a Max Binary Heap:



When an element is pushed to the priority queue, it is simply added to the heap.



When an element is popped from the priority queue

Priority Queue - Representation on a binary heap



When representing a Priority Queue using a Max Binary Heap:



When an element is pushed to the priority queue, it is simply added to the heap.



When an element is popped from the priority queue, the root is removed from the heap.



Top

Priority Queue - Representation on a binary heap



When representing a Priority Queue using a Max Binary Heap:



When an element is pushed to the priority queue, it is simply added to the heap.



When an element is popped from the priority queue, the root is removed from the heap.



Top simply returns the pair in the root of the heap.

Representing a Priority Queue using a Max-Heap



A Priority Queue element's representation:

PQElement:

Representing a Priority Queue using a Max-Heap



A Priority Queue element's representation:

PQElement:

e: TElem

p: TPriority



Priority Queue representation using a Max-Heap:

PriorityQueue:

Representing a Priority Queue using a Max-Heap



A Priority Queue element's representation:

PQElement:

e: TElem

p: TPriority



Priority Queue representation using a Max-Heap:

PriorityQueue:

capacity: Integer

length: Integer

elems: PQElement[]

R: TPriority \times TPriority \rightarrow {True, False}

Priority Queue represented on Heap - init



Creating an empty PQ:

subalgorithm `init(pq, R)` **is:**

Priority Queue represented on Heap - init



Creating an empty PQ:

subalgorithm init(pq, R) **is:**

//R\$ TPriority \times TPriority \rightarrow {True, False}

pq.R \leftarrow R

pq.length \leftarrow 0

pq.capacity \leftarrow @ an initial capacity

@ allocate pq.elems of capacity pq.capacity

end-subalgorithm

Priority Queue represented on Heap - add



Inserting a new element into the PQ:

subalgorithm push(pq, e, p) **is:**

//pq - a pq

//e - the element to be added

//p - the priority of the element to be added

if pq.length = pq.capacity **then**

 @ resize

end-if

pq.length \leftarrow pq.length + 1

pq.elems[pq.length].e \leftarrow e

pq.elems[pq.length].p \leftarrow p

bubble-up(pq, pq.length)

end-subalgorithm

Priority Queue represented on Heap - push



Bubble-up for restoring the heap property:

subalgorithm bubble-up (pq, p) **is:**

//pq - a Priority Queue

//p - position from which we bubble the new node up

pos \leftarrow p

elem \leftarrow pq.elems[p]

parent \leftarrow [p / 2]

while pos > 1 **and not**(pq.R(pq.elems[parent].p, elem.p)) **execute**

//move parent down

pq.elems[pos] \leftarrow pq.elems[parent]

pos \leftarrow parent

parent \leftarrow [pos/2]

end-while

pq.elems[pos] \leftarrow elem

end-subalgorithm

Priority Queue represented on Heap - push - complexity



What is the time complexity of *push*?

Priority Queue represented on Heap - push - complexity



What is the time complexity of *push*?



$O(\log_2 n)$ amortized

Priority Queue represented on Heap - pop



Popping from a Priority Queue:

function pop(pq) **is:**

//pq - is a pq

if pq.length = 0 **then**

 @ error - empty Priority Queue

end-if

deletedElem \leftarrow (pq.elems[1].e, pq.elems[1].p)

pq.elems[1] \leftarrow pq.elems[pq.length]

pq.length \leftarrow pq.length - 1

bubble-down(pq, 1)

remove \leftarrow deletedElem

end-function



Bubble-down for restoring the heap property:

subalgorithm bubble-down(pq, p) **is:**

pos \leftarrow p

elem \leftarrow pq.elems[p]

while pos \leq [pq.length/2] **execute**

maxChild \leftarrow pos*2

if pos*2+1 \leq pq.length **and** pq.R(pq.elems[2*pos+1].p, pq.elems[2*pos].p)

then *//it has two children and the right is greater*

maxChild \leftarrow pos*2 + 1

end-if

if pq.R(pq.elems[maxChild], elem) **then**

tmp \leftarrow pq.elems[pos]

pq.elems[pos] \leftarrow pq.elems[maxChild]

pq.elems[maxChild] \leftarrow tmp

pos \leftarrow maxChild

else

pos \leftarrow pq.length + 1 *//to stop the while loop*

end-if

end-while

end-subalgorithm

Priority Queue represented on Heap - pop - complexity



What is the time complexity of *pop*?

Priority Queue represented on Heap - pop - complexity



What is the time complexity of *pop*?



$O(\log_2 n)$

Priority Queue represented on Heap - top



Popping from a Priority Queue:

function top(pq) **is:**

//pq - is a pq

Priority Queue represented on Heap - top



Popping from a Priority Queue:

function top(pq) **is:**

//pq - is a pq

if pq.length = 0 **then**

 @ error - empty Priority Queue

end-if

top \leftarrow (pq.elems[1].e, pq.elems[1].p)

end-function



Complexity:

Priority Queue represented on Heap - top



Popping from a Priority Queue:

function top(pq) **is:**

//pq - is a pq

if pq.length = 0 **then**

 @ error - empty Priority Queue

end-if

top \leftarrow (pq.elems[1].e, pq.elems[1].p)

end-function



Complexity: $\Theta(1)$

Priority Queue represented on Heap - isEmpty



Checking if a Priority Queue is empty:

function isEmpty(pq) **is:**

//pq - is a pq

if pq.length = 0 **then**

isEmpty \leftarrow True

else

isEmpty \leftarrow False

end-if

end-function



Complexity: $\Theta(1)$

Priority Queue in programming languages



ADT Priority Queue in programming languages:

- PriorityQueue in Java
 - represented using the Heap data structure
- PriorityQueue in Python (queue module)
 - represented using the Heap data structure
- priority_queue in C++ STL
 - represented using the Heap data structure

Priority Queue - Applications



Applications of priority queues:



Graph algorithms

- Implementation of Priority Queues used in Dijkstra's algorithm (shortest path) or Prim's algorithm (minimum spanning tree)



Data compression

- Huffman coding



Resources allocation

- To efficiently allocate resources in a system, such as memory blocks or CPU time, by processing requests in order of priority



Artificial Intelligence

- A* search algorithm



Operating Systems

- In the load balancing algorithms

Heap-sort - Naive approach




The initial sequence: [6, 1, 3, 9, 11, 4, 2, 5]

Heap-sort - Naive approach

► The initial sequence: [6, 1, 3, 9, 11, 4, 2, 5]


► If we **add all elements into a min-heap**, we get:



Heap-sort - Naive approach

► The initial sequence: [6, 1, 3, 9, 11, 4, 2, 5]

► If we **add all elements into a min-heap**, we get:



► Then, if we **remove all the elements**, one-by-one, we obtain:
1, 2, 3, 4, 5, 6, 9, 11.

Heap-sort - Naive approach



What is the time complexity of the heap-sort algorithm previously described?

Heap-sort - Naive approach



What is the time complexity of the heap-sort algorithm previously described?



$O(n \cdot \log_2 n)$

Heap-sort - Naive approach



What is the time complexity of the heap-sort algorithm previously described?



$O(n \cdot \log_2 n)$



What is the extra space complexity of the heap-sort algorithm previously described?

Heap-sort - Naive approach



What is the time complexity of the heap-sort algorithm previously described?



$O(n \cdot \log_2 n)$



What is the extra space complexity of the heap-sort algorithm previously described?



$\Theta(n)$

Heap-sort - Better approach

- ▶ We start by transforming the unsorted array into a max-heap.
- ▶ The second half of the array contain leaves, so they can be left where they are.
- ▶ Starting from the last non-leaf element (and going towards the beginning of the array), we just call *bubble-down* for every element.

Heap-sort - Better approach



Heap-sort - better approach:

function build-max-heap(a, n) **is:**

//a - an array of length n

heap.elems \leftarrow a

heap.length \leftarrow n

heap.capacity \leftarrow n

for i \leftarrow [n/2], 1, -1 **execute**

 bubble-down(heap, i)

end-for

build-max-heap \leftarrow heap

end-function



Time complexity: $O(n)$

Heap-sort - Better approach

- After transforming the unsorted array into a max-heap:
 - ▶ The maximum element is stored in the root, so at index 1 \Rightarrow we swap it with the one at the last index
 - ▶ We discard the last element in the heap, by decrementing the length of the heap
 - ▶ The root element is the only that may violate the heap property \Rightarrow we *bubble-down* it
 - 🕒 We repeat the process until the length of the heap is 1

Heap-sort - Better approach



Heap-sort - better approach:

```
subalgorithm heapsort(a, n) is:  
//a - an array of lenght n  
  heap  $\leftarrow$  build-max-heap(a, n)  
  for i  $\leftarrow$  n, 2, -1 execute  
    aux  $\leftarrow$  a[i]  
    a[i]  $\leftarrow$  a[1]  
    a[1]  $\leftarrow$  aux  
    heap.length  $\leftarrow$  heap.length-1  
    bubbleDown(a, 1)  
  end-for  
end-subalgorithm
```

Heap-sort - Complexity



Time complexity of this heap-sort is $O(n \cdot \log_2 n)$.



build-max-heap runs in $O(n)$.



bubble-down runs in $O(\log_2 n)$ and we call it $n - 1$ times.



Extra-space complexity of this approach is $\Theta(1)$.



Bibliography

- David M. Mount, *Lecture notes for the course Data Structures* (CMSC 420), at the Dept. of Computer Science, University of Maryland, College Park, 2001
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein, *Introduction to Algorithms*, Third Edition, The MIT Press, 2009
- Narasimha Karumanchi, *Data Structures and Algorithms Made Easy: Data Structures and Algorithmic Puzzles*, Fifth Edition, 2016
- Clifford A. Shaffer, *A Practical Introduction to Data Structures and Algorithm Analysis*, Third Edition, 2010

► THANK
E
A
P

Y P
R I Q
O U
R E
I U
T E
Y