

Lecture 12

Query Optimization in Relational Databases. Evaluating Relational Algebra Operators.

SQL Statement Execution

- Client application – the SQL statement execution request - should be performed for any query and provide a minimum response time
- Stages for the statement execution:
 - client: generate SQL statement (non-procedural language) and send it to server
 - server:
 - analyze SQL statement (in syntactically manner)
 - translate statement into an internal form (by using relational algebra expression)
 - transform internal form into an optimal form (similar to relational algebra)
 - generate a procedural execution plan
 - evaluate procedural plan, send result to client

SQL Statement Execution

- The operators that are used in the query process are:

- **Selection:** $\sigma_C(R)$
- **Projection:** $\pi_\alpha(R)$
- **Cross-product:** $R_1 \times R_2$
- **Union:** $R_1 \cup R_2$
- **Intersection:** $R_1 \cap R_2$
- **Set-difference:** $R_1 - R_2$
- **Theta join:** $R_1 \otimes_\theta R_2$
- **Natural join:** $R_1 * R_2$

- **Left outer join:** $R_1 \bowtie_C R_2$
- **Right outer join:** $R_1 \bowtie_C R_2$
- **Full outer join:** $R_1 \bowtie_C R_2$
- **Left semi join:** $R_1 \triangleright R_2$
- **Right semi join:** $R_1 \triangleleft R_2$
- **Division:** $R_1 \div R_2$
- **Duplicate elimination:** $\delta(R)$
- **Sorting:** $S_{\{list_attributes\}}(R)$
- **Grouping:** $\gamma_{\{list1_attributes\}} GROUP BY \{list2_attributes\} (R)$

SQL Statement Execution

- An SQL query may be written in multiple ways

Example: Consider the following relational database

Specialization[SpId, SpecName, SpecDescription, NoOfStudents]

Group[Gid, NoOfStudents, StudyYear, *SpId*]

Student[Sid, FirstName, LastName, Age, Email, *Gid*]

SQL query: display the students (FirstName, LastName, Age, StudyYear, SpecName) from a given specialization (e.g. SpId=3 can be a parameter) with the Age>=20 (can be a parameter)

* The provided solutions are equivalent (or, provide the same answer)

Solution 1:

SELECT FirstName, LastName, Age, StudyYear, SpecName

FROM Student S, Group G, Specialization Sp

WHERE S.Gid=G.Gid AND G.SpId=Sp.SpId AND Sp.SpId=3 AND Age>=20

$$\pi_{\text{FirstName, LastName, Age, StudyYear, SpecName}} \left(\sigma_{\text{Student.Gid=Group.Gid AND Group.SpId=Specialization.SpId AND Sp.SpId=3 AND Age} \geq 20} \text{Student} \times \text{Group} \times \text{Specialization} \right)$$

SQL Statement Execution

Specialization[SpId, SpecName, SpecDescription, NoOfStudents]

Group[Gid, NoOfStudents, StudyYear, *SpId*]

Student[Sid, FirstName, LastName, Age, Email, *Gid*]

SQL query: display the students (FirstName, LastName, Age, StudyYear, SpecName) from a given specialization (e.g. SpId=3 can be a parameter) with the Age \geq 20 (can be a parameter)

Solution 2:

```
SELECT FirstName, LastName, Age, StudyYear, SpecName
FROM (Student S INNER JOIN Group G ON S.Gid=G.Gid)
INNER JOIN Specialization Sp ON G.SpId=Sp.SpId
WHERE Sp.SpId=3 AND Age $\geq$ 20
```

$$\pi_{\text{FirstName, LastName, Age, StudyYear, SpecName}}(\sigma_{\text{Specialization.SpId}=3 \text{ AND Age} \geq 20}(\text{Student} \otimes_{\text{Student.Gid=Group.Gid}} \text{Group}) \otimes_{\text{Specialization.SpId=Group.SpId}} \text{Specialization}))$$

SQL Statement Execution

Specialization[SpId, SpecName, SpecDescription, NoOfStudents]

Group[Gid, NoOfStudents, StudyYear, *SpId*]

Student[Sid, FirstName, LastName, Age, Email, *Gid*]

SQL query: display the students (FirstName, LastName, Age, StudyYear, SpecName) from a given specialization (e.g. SpId=3 can be a parameter) with the Age>=20 (can be a parameter) * *that are assign in the given group (e.g. Gid=3 can be a parameter)*

Solution 3:

```
SELECT FirstName, LastName, Age, StudyYear, SpecName
```

```
FROM ( (SELECT FirstName, LastName, Age, Gid
```

```
FROM Student
```

```
WHERE Age>=20) S
```

```
INNER JOIN (SELECT * FROM Group WHERE SpId=3) G ON S.Gid=G.Gid
```

```
) INNER JOIN
```

```
(SELECT SpId, SpecName FROM Specialization WHERE SpId=3) Sp ON G.SpId=Sp.SpId
```

$$\pi_{FirstName, LastName, Age, StudyYear, SpecName} \left(\left(\pi_{FistName, LastName, Age, Gid} \left(\left(\sigma_{Age \geq 20}(Student) \right) \right) \otimes_{Gid_condition} \left(\sigma_{SpId=3}(Group) \right) \right) \otimes_{SpId_condition} \left(\pi_{SpId, SpecName} \left(\left(\sigma_{SpId=3}(Specialization) \right) \right) \right)$$

- An evaluation tree can be constructed for a relational algebra expression
- Problems:
 - Which version is better?
 - When generating the execution plan:
 - Which parameters are optimized?
 - What information is required?
 - What can the Optimizer (a component of the DBMS) do?

Relational Algebra Operators - Evaluation

- Operands for relational operators:
 - database tables (can have attached indexes)
 - temporary tables (obtained by evaluating some relational operators)
- several evaluation algorithms could be used for a relational algebra operator
- when generating the execution plan:
 - choose the algorithm with the lowest complexity (for the current database context); take into account data from the system catalog, statistical information

Relational Algebra Operators - Evaluation

Algorithms

1. Table Scan

- A lot of operators require a full scan of the entire table
- Let b_R be the number of blocks storing a table's records
 - sequential search algorithm - approximately $b_R/2$ blocks are necessary (on average) when performing a sequential search on a key value
 - all blocks must be brought into main memory when performing a sequential search on a non-key field
- High transfer time for large tables

2. Index Seek

- Searching for a key value K_0 by having a condition of the form: $K = K_0$
 - Search can be performed
 - explicit (searching a table, evaluating a join)
 - implicit (checking a key constraint)
 - Examine an index (stored as a B-tree, B+ tree) created:
 - via a key constraint
 - with the CREATE INDEX statement
- where K can be a simple or composite key

Relational Algebra Operators - Evaluation

Algorithms

3. Index Scan

- Evaluating $\sigma_C(R)$, where condition C has the form:
 - $A < v, A \leq v, A > v, A \geq v, A \text{ IS [NOT] NULL}$ – index built for a key A
 - $A = v, A < v, A \leq v, A > v, A \geq v, A \text{ IS [NOT] NULL}$ – index built for a non-key field A
- From the partial / total index scan are obtained desired records' addresses
- Get records from the relation; some blocks can be read multiple times

- A **join** can be defined as a cross-product followed by a selection
- Joins arise more often in practice than cross-products
- In general, the result of a cross-product is much larger than the result of a join
- It is important to implement the join without materializing the underlying cross-product, by applying selections and projections as soon as possible, and materializing only the subset of the cross-product that will appear in the result of the join

Relational Algebra Operators - Evaluation

Cross Join

- The algorithm is used to evaluate a cross-product:
 - $R \text{ CROSS JOIN } S$
 - $R \text{ INNER JOIN } S \text{ ON } C$ (C evaluates to TRUE)
 - $\text{SELECT ... FROM } R, S \text{ ...}$, no join condition between R and S
- b_R, b_S – the number of blocks storing R and S
- m, n – the number of blocks from R and S that can simultaneously appear in the main memory (there are $m+n$ buffers for the 2 tables)
- Algorithm used to generate the cross-product $\{(r, s) \mid r \in R, s \in S\}$:
 - for every group of max. m blocks in R :
 - read the group of blocks from R into main memory;
 - let M_1 be the set of records in these blocks
 - for every group of max. n blocks in S :
 - read the group of blocks from S into main memory;
 - let M_2 be the set of records in these blocks
 - for every $r \in M_1$:
 - for every $s \in M_2$:
 - add (r, s) to the result

Relational Algebra Operators - Evaluation

Cross Join

- Algorithm complexity: total number of read blocks (from the 2 tables):

$$b_R + [b_R/m] * b_S$$

(number of blocks in R; for every group of max. m blocks in R, read S)

- To minimize this value, m should be maximized (the other operands are constants); one buffer can be used for S (so $n = 1$), while the remaining space can be used for R (m max.)
- Switch the 2 relations (in the algorithm and when computing the complexity) \Rightarrow complexity:

$$b_S + [b_S/n] * b_R$$

- Choose better version
- If $b_R \leq m$ or $b_S \leq n \Rightarrow$ complexity $b_R + b_S$

Nested Loops Join

- The Cross Join algorithm can be used to evaluate a join between 2 tables
- For every element (r, s) in the cross-product, evaluate the condition in the join operator
- The elements (r, s) that don't meet the join condition are eliminated

Relational Algebra Operators - Evaluation

Indexed Nested Loops Join

- The algorithm is used to evaluate $R \bowtie_C S$, where $C \equiv (R.A=S.B)$, and there is an index on A (in R) or on B (in S)
- For the algorithm, it is assumed that there is an index on column B in table S
 - for every block in R:
 - read the block into main memory
 - let M be the set of records in the block
 - for every $r \in M$:
 - determine $v = \pi_A(r)$
 - use the index on B in S to determine records $s \in S$ with value v for B
 - for every such record s, the pair (r,s) is added to the result

Depending on the type of index, there can be at most 1 / multiple matching records in S

Relational Algebra Operators - Evaluation



Merge Join

- The algorithm is used to evaluate $R \bowtie_C S$, where $C \equiv (R.A=S.B)$, and there are no indexes on A (in R) and B (in S)
 - sort R and S on the columns used in the join: R on A, S on B
 - scan obtained tables
 - let r in R and s in S be 2 current records
 - if $r.A = s.B$
 - add (r', s') to the result; r' is in the set of all consecutive records in R with $A = r.A$
 - add (r', s') to the result; s' is in the set of all consecutive records in S with $B = s.B$
 - $\text{next}(r)$
 - $\text{next}(s)$ (get a record with the next value for A and B)
 - if $r.A < s.B$
 - $\text{next}(r)$ (determine record in sorted R with the next value for A)
 - if $r.A > s.B$
 - $\text{next}(s)$ (determine record in sorted S with the next value for B)

Relational Algebra Operators - Evaluation

Merge Join

- The algorithm replaces an evaluation tree with another evaluation tree



Relational Algebra Operators - Evaluation

Hash Join

- The algorithm is used to evaluate $R \bowtie_C S$, where $C \equiv (R.A = S.B)$
 - a. partitioning phase – hash R and S on the join column, use the same hash function $h \Rightarrow$ partitions
 - b. probing phase – tuples in partition R_x are compared only with tuples in partition S_x (tuples in partition R_1 cannot join with tuples in partition S_2 , for instance, as they have a different hash value)

Outer Joins

- The condition join algorithms has to be adapted

Operations on Sets of Records: $R \cup S$, $R - S$, $R \cap S$

- The previous algorithms should be adapted
- For intersection:
 - sort R using all columns
 - sort S using all columns
 - scan sorted R and S , write in the result only the tuples in R that also appear in S

Relational Algebra Equivalences

- SQL statement - transformed into a relational algebra expression (based on a set of transformation rules for the clauses that appear in the statement)
- transform relational expression (such that the evaluation algorithm has a lower complexity)
- certain transformation rules are used (mathematical properties of the relational operators)

- $\sigma_C(\pi_\alpha(R)) = \pi_\alpha(\sigma_C(R))$

- selection reduces the number of records for projection
- in the second expression, the projection operator analyzes fewer records
- optimization - algorithm that evaluates both operators in a single pass of R

- perform one pass instead of 2:

$$\sigma_{C1}(\sigma_{C2}(R)) = \sigma_{C1 \text{ AND } C2}(R)$$

- replace cross-product and selection by condition join (a number of condition join algorithms don't evaluate the cross-product):

$$\sigma_C(R \times S) = R \Join_C S$$

where C - join condition between R and S

Relational Algebra Equivalences

- Let R and S be compatible schemas

$$\sigma_C(R \cup S) = \sigma_C(R) \cup \sigma_C(S)$$

$$\sigma_C(R \cap S) = \sigma_C(R) \cap \sigma_C(S)$$

$$\sigma_C(R - S) = \sigma_C(R) - \sigma_C(S)$$

- $\sigma_C(R \times S)$ with the particular cases:

- C contains only attributes from R:

$$\sigma_C(R \times S) = \sigma_C(R) \times S$$

- $C = C1 \text{ AND } C2$, C1 contains only attributes from R, C2 - only attributes from S:

$$\sigma_{C1 \text{ AND } C2}(R \times S) = \sigma_{C1}(R) \times \sigma_{C2}(S)$$

- $C = C1 \text{ AND } C2$, C2 - join condition between R and S:

$$\sigma_{C1 \text{ AND } C2}(R \times S) = \sigma_{C1}(R \bowtie_{C2} S)$$

- $\pi_\alpha(R \cup S) = \pi_\alpha(R) \cup \pi_\alpha(S)$

- $\pi_\alpha(R \bowtie_C S) = \pi_\alpha(\pi_{\alpha1}(R) \bowtie_C \pi_{\alpha2}(S))$

where, $\alpha1$ are attributes in R that appear in α or C and $\alpha2$ are attributes in S that appear in α or C

Relational Algebra Equivalences

- associativity and commutativity for some relational operators
 - associativity and commutativity for \cup and \cap
 - associativity for the cross-product and the natural join
 - "equivalent" results (same records, but different column order) when commuting operands in \times and certain join operators
 - $R \times S = S \times R$ – when using the Cross Join algorithm, the order of the data sources is important
- transitivity of some relational operators for the join operators - additional filters could be applied before the join:
 - $(A > B \text{ AND } B > 5) \equiv (A > B \text{ AND } B > 5 \text{ AND } A > 5)$

Example: A is in R, B is in S: $R \otimes_{A > B \text{ AND } B > 5} S = (\sigma_{A > 5}(R)) \otimes_{A > B} (\sigma_{B > 5}(S))$

- $(A = B \text{ AND } B = 5) \equiv (A = B \text{ AND } B = 5 \text{ AND } A = 5)$

Example: A is in R, B is in S: $R \otimes_{A = B \text{ AND } B = 5} S = (\sigma_{A = 5}(R)) \otimes_{A = B} (\sigma_{B = 5}(S))$

- evaluating $\sigma_C(R)$, where $C \equiv (R. A \in \delta(\pi_{\{B\}}(S)))$; avoid evaluating C for every record of R; the initial evaluation is equivalent to: $R \otimes_{R.A = S.B} (\delta(\pi_{\{B\}}(S)))$

Relational Algebra Equivalences

Example: Consider the (same) relational database

Specialization[SpId, SpecName, SpecDescription, NoOfStudents]

Group[Gid, NoOfStudents, StudyYear, *SpId*]

Student[Sid, FirstName, LastName, Age, Email, *Gid*]

SQL query: display the students (FirstName, LastName, Age, StudyYear, SpecName) from a given specialization (e.g. SpId=3 can be a parameter) with the Age>=20 (can be a parameter)

SELECT FirstName, LastName, Age, StudyYear, SpecName

FROM Student S, Group G, Specialization Sp

WHERE S.Gid=G.Gid AND G.SpId=Sp.SpId AND Sp.SpId=3 AND Age>=20

- Denote by $C \equiv S.Gid=G.Gid \text{ AND } G.SpId=Sp.SpId \text{ AND } Sp.SpId=3 \text{ AND } Age \geq 20$
- Denote by $\beta = \{FirstName, LastName, Age, StudyYear, SpecName\}$ – the attributes from the SELECT clause
- So, the relational expression is

$$\pi_{\beta}(\sigma_C (Student \times Group \times Specialization))$$

Relational Algebra Equivalences

Pay attention to the following transformations, by using the previously discussed rules:

- Associativity for \times
 - $Student \times Group \times Specialization = (Student \times Group) \times Specialization$
 - $Student \times Group \times Specialization = Student \times (Group \times Specialization)$
- Commute σ with \times (use a particular case); the transitivity of the equality operator will be
 - $(G.SpId=Sp.SpId \text{ AND } Sp.SpId=3) \equiv (G.SpId=Sp.SpId \text{ AND } Sp.SpId=3 \text{ AND } G.SpId=3)$
- For

S.Gid=G.Gid AND G.SpId=Sp.SpId AND Sp.SpId=3 AND Age>=20 AND G.SpId=3

C1

C2

C3

C4

C5

$$\sigma_C(Student \times Group \times Specialization) = \sigma_{C1 \text{ AND } C2} \left((\sigma_{C4}(Student) \times \sigma_{C5}(Group)) \times \sigma_{C3}(Specialization) \right)$$

or

$$\sigma_C(Student \times Group \times Specialization) = \sigma_{C1 \text{ AND } C2} (\sigma_{C4}(Student)) \times (\sigma_{C5}(Group) \times \sigma_{C3}(Specialization))$$

- Replace selection and cross-product with condition join

$$\sigma_C(Student \times Group \times Specialization) = ((\sigma_{C4}(Student)) \otimes_{C1} (\sigma_{C5}(Group))) \otimes_{C2} (\sigma_{C3}(Specialization))$$

or

$$\sigma_C(Student \times Group \times Specialization) = ((\sigma_{C4}(Student))) \otimes_{C1} ((\sigma_{C5}(Group)) \otimes_{C2} (\sigma_{C3}(Specialization)))$$

Relational Algebra Equivalences

- Choose a version based on statistical information from the database; next it is considered the first version:

$$e = \pi_{\beta}(((\sigma_{C4}(Student)) \otimes_{C1} (\sigma_{C5}(Group))) \otimes_{C2} (\sigma_{C3}(Specialization)))$$

- commute π with join

- $\beta1 = \{FirstName, LastName, Age, Gid\}$ - useful for β and join
- $\beta2 = \{Gid, StudyYear, SpId\}$ - useful for β and join
- $\beta3 = \{SpId, SpecName\}$ - useful for β and join

$$e = \pi_{\beta}(((\pi_{\beta1}(\sigma_{C4}(Student)) \otimes_{C1} (\pi_{\beta2}(\sigma_{C5}(Group)))) \otimes_{C2} (\pi_{\beta3}(\sigma_{C3}(Specialization))))$$

- This expression correspond to *Solution 3*:

```
SELECT FirstName, LastName, Age, StudyYear, SpecName
FROM (
```

```
  (SELECT FirstName, LastName, Age, Gid
```

```
  FROM Student
```

```
  WHERE Age >= 20) S
```

```
  INNER JOIN (SELECT * FROM Group WHERE Gid=3) G ON S.Gid=G.Gid
```

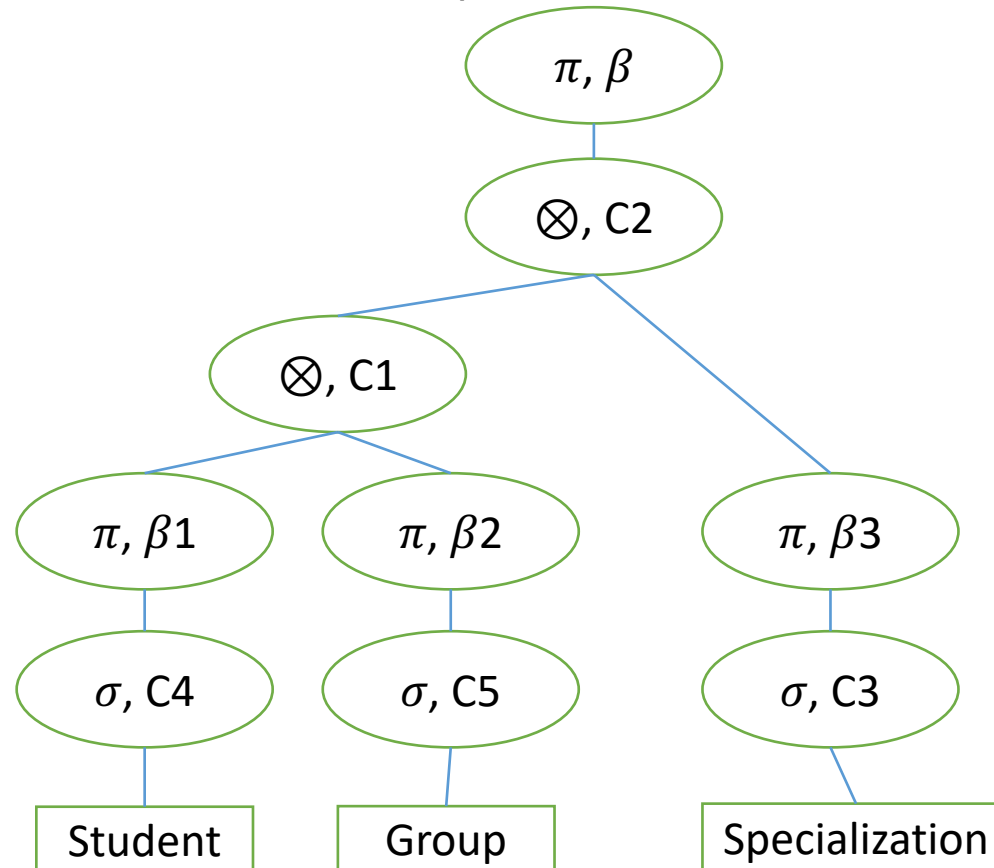
```
  ) INNER JOIN
```

```
  (SELECT SpId, SpecName FROM Specialization WHERE SpId=3) Sp ON G.SpId=Sp.SpId
```

Relational Algebra Equivalences

- An evaluation tree can be constructed for the last version of the relational algebra expression
- Using information from the system catalog and possibly statistical information, an execution plan can be generated from the last version of the expression; every relational operator is replaced by an evaluation algorithm

$$e = \pi_{\beta}(((\pi_{\beta_1}(\sigma_{C_4}(Student)) \otimes_{C_1}(\pi_{\beta_2}(\sigma_{C_5}(Group)))) \otimes_{C_2}(\pi_{\beta_3}(\sigma_{C_3}(Specialization))))))$$



References:

- C.J. Date, *An Introduction to Database Systems (8th Edition)*, Addison-Wesley, 2003.
- H. Garcia-Molina, J. Ullman, J. Widom, *Database Systems: The Complete Book*, Prentice Hall Press, 2008.
- G. Hansen, J. Hansen, *Database Management And Design (2nd Edition)*, Prentice Hall, 1996.
- R. Ramakrishnan, J. Gehrke, *Database Management Systems*, McGraw- Hill, 2007.
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- R. Ramakrishnan, J. Gehrke, *Database Management Systems (2nd Edition)*, McGraw-Hill, 2000.
- A. Silberschatz, H. Korth, S. Sudarshan, *Database System Concepts*, McGraw-Hill, 2010.
<http://codex.cs.yale.edu/avi/db-book/>
- L. Țâmbulea, *Curs Baze de date*, Facultatea de Matematică și Informatică, UBB, 2013-2014.
- J. Ullman, J. Widom, *A First Course in Database Systems*,
<http://infolab.stanford.edu/~ullman/fcdb.html>