

Lecture 14

# DATABASES

Transactions. Concurrency  
Control

# Concurrency in a DBMS

## **Scenarios**

### 1. Train reservation system

- Travel agent 1 (TA1) processes the request of customer 1 (C1) related to a seat on the train 10 → system retrieves from the database the tuple <train: 10, available seats: 1> and displays it to TA1
- C1 takes time to decide whether or not to make the reservation
- Travel agent 2 (TA2) processes the request of customer 2 (C2) related to a seat on the train 10 → system retrieves from the database the tuple <train: 10, available seats: 1> and displays it to TA2
- C2 immediately reserves the seat
- system:
  - update the tuple retrieved by TA2: <train: 10, available seats: 0>
  - updates the database, i.e. replaces flight 10's tuple in the database with the new tuple <flight: 10, available seats: 0>
- C1 also decides to reserve the seat
- however, TA1 is still working on the previously retrieved tuple, not knowing another user has changed the data

# Concurrency in a DBMS

## ***Scenarios***

- system:
  - update the tuple retrieved by TA1: <train: 10, available seats: 0>
  - updates the database, i.e. replaces flight 10's tuple in the database with the new tuple <flight: 10, available seats: 0>
- The last remaining seat on train 10 has been sold 2 times!

## 2. Bank system

- 500 accounts:  $A_1, A_2, \dots, A_{500}$
- program P1 is computing the total balance:  $A_1 + A_2 + \dots + A_{500}$
- program P2 transfers 1000 lei from  $A_1$  to  $A_{500}$
- if P2's transfer occurs after P1 has “seen”  $A_1$ , but before it has “seen”  $A_{500}$  => the total balance will be incorrect:  $A_1 + A_2 + \dots + A_{500} + 500$
- The objective is to prevent such anomalies without disallowing concurrent access!

# Concurrency in a DBMS

## ***Transactions***

- The database is partitioned into *items* (e.g. tuple, attribute)

## Transaction

- any *one execution* of a program in a DBMS
- sequence of one or more operations on a database: *read / write / commit / abort*
- *commit* occurs if and only if the transaction is not aborted
- final operation: commit or abort

Let *I* be an item in a DB (database)

- fundamental operations in a transaction *T* (*T* can be omitted when clear from context):
  - *read*(*T*, *I*)
  - *write*(*T*, *I*)
  - *commit*(*T*) - successful completion: *T*'s changes are to become permanent
  - *abort*(*T*) - *T* is *rolled back*: its changes are undone

# Concurrency in a DBMS

## Transaction properties - ACID

- properties that a DBMS must ensure for a transaction in order to maintain data in the presence of concurrent execution / system failures

### ***Atomicity***

- a transaction is atomic
- either all the operations in the transaction are executed, or none are (*all-or-nothing*)

### ***Consistency***

- a transaction must preserve the consistency of the database after execution, i.e. a transaction is a correct program

### ***Isolation***

- a transaction is protected from the effects of concurrently scheduling other transactions

### ***Durability***

- committed changes are persisted to the database
- the effects of a committed transaction should persist even if a system crash occurs before all the changes have been written to disk

# Concurrency in a DBMS

## **Atomicity**

- a transaction can commit / abort:
  - *commit* – after finalizing all its operations (successful completion)
  - *abort* – after executing some of its operations; the DBMS can itself abort a transaction
- user's perspective:
  - either all operations in a transaction are executed, or none are
  - the transaction is treated as an indivisible unit of execution
- the DBMS *logs* all transactions' actions, so it can *undo* them if necessary (i.e. undo the actions of incomplete transactions)

e.g. consider the following transaction transferring money from Account\_A to Account\_B:

Account\_A  $\leftarrow$  Account\_A - 100


Account\_B  $\leftarrow$  Account\_B + 100

- if the transaction fails in the middle, 100 lei must be put back into Account\_A, i.e. partial effects are undone

# Concurrency in a DBMS

## **Consistency**

- in the absence of other transactions, a transaction that executes on a consistent database state leaves the database in a consistent state



- transactions do not violate the integrity constraints (ICs) specified on the database (e.g. restrictions declared via CREATE TABLE statements) and enforced by the DBMS
- however, users can have consistency criteria that go beyond the expressible integrity constraints

# Concurrency in a DBMS

## **Consistency**

e.g. consider again the money transfer transaction:

$\text{Account\_A} \leftarrow \text{Account\_A} - 100$

$\text{Account\_B} \leftarrow \text{Account\_B} + 100$

- assume the transaction starts execution at time  $t_i$  and completes at  $t_f$
- consistency constraint:  $\text{Total}(\text{Account\_A}, \text{Account\_B}, t_i) = \text{Total}(\text{Account\_A}, \text{Account\_B}, t_f)$
- it is the responsibility of the user to write a correct transaction that meets the above constraint
- the DBMS does not understand the semantics of the data, of user actions (e.g. how is the interest computed, the fact that Account\_B must be credited with the amount of money debited from Account\_A)



# Concurrency in a DBMS

## ***Isolation***

- transactions are protected from the effects of concurrently scheduling other transactions
- users concurrently submit transactions to the DBMS, but can think of their transactions as if they were executing in isolation, independently of other users' transactions
- i.e. users must not deal with arbitrary effects produced by other concurrently running transactions; a transaction is seen as if it were in single-user mode

## ***Durability***

- once a transaction commits, the system must guarantee the effects of its operations won't be lost, even if failures subsequently occur
- i.e. once the DBMS informs the user a transaction has committed, its effects should persist, even if a system crash occurs before all changes have been saved on the disk

# Concurrency in a DBMS

## ***Lock-Based Concurrency Control***

- *lock*
  - a tool used by the transaction manager to control concurrent access to data
  - prevents a transaction from accessing a data object while another transaction is accessing the object
- *transaction protocol*
  - a set of rules enforced by the transaction manager and obeyed by all transactions
  - e.g. – simple protocol: before a transaction can read / write an object, it must acquire an appropriate lock on the object
  - locks in conjunction with transaction protocols allow interleaved executions
- *SLock (shared or read lock)*
  - if a transaction holds an SLock on an object, it can read the object, but it cannot modify it
- *XLock (exclusive or write lock)*
  - if a transaction holds an XLock on an object, it can both read and write the object

# Concurrency in a DBMS

## ***Lock-Based Concurrency Control***

- if a transaction holds an SLock on an object, other transactions can be granted SLocks on the object, but they cannot acquire XLocks on it
- if a transaction holds an XLock on an object, other transactions can be granted either SLocks or Xlock on the object

	Shared	Exclusive
Shared	Yes	No
Exclusive	No	No

# Concurrency in a DBMS

## ***Transaction Support in SQL - Isolation Levels***

- SQL provides support for users to specify various aspects related to transactions e.g. isolation levels
- *isolation level*
  - a transaction's characteristic
  - determines the degree to which a transaction is isolated from the changes made by other concurrently running transactions
- greater concurrency -> concurrency anomalies
- 4 isolation levels
  - READ UNCOMMITTED
  - READ COMMITTED
  - REPEATABLE READ
  - SERIALIZABLE

# Concurrency in a DBMS

## ***Transaction Support in SQL - Isolation Levels***

- isolation levels can be set with the following command:  
SET TRANSACTION ISOLATION LEVEL *isolevel*  
*e.g.* SET TRANSACTION ISOLATION LEVEL REPEATABLE READ

Examples on Article(Aid, Title, PrincipalAuthor, Year, Type)

- Article contains (at least) the following tuples:  
(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

### READ UNCOMMITTED

- no S locks when reading data
- a transaction can see uncommitted changes carried out by another transaction

# Concurrency in a DBMS

## Dirty reads

- a transaction reads data that has been modified by another transaction, but not committed
  - T2 reads *John Paul* – under READ UNCOMMITTED
- (3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

UPDATE Article  
SET PrincipalAuthor='John Paul'  
WHERE Aid=3

T1 rolls back

T2

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3

T2 commits

time

# Concurrency in a DBMS

## *Isolation Levels*

### READ COMMITTED

- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object (i.e. the last transaction that modified the object is complete)
- exclusive locks are released at the end of the transaction
- shared locks are immediately released
- a transaction cannot see uncommitted changes carried out by another transaction
- a transaction can change data that has been read by the current transaction, prior to its completion
- it is the default isolation level in SQL SERVER

# Concurrency in a DBMS

## Dirty reads

- T2 reads *Popescu Dan*– under READ COMMITTED

(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

UPDATE Article  
SET PrincipalAuthor='John Paul'  
WHERE Aid=3

T1 rolls back

T2

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3

T2 commits

time



# Concurrency in a DBMS

## Non-repeatable reads

- T2 under READ COMMITTED; T1 can change data read by T2 while T2 is still in progress
- T2 reads 2 different values: *Popescu Dan* and *John Paul* – under READ COMMITTED

(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

UPDATE Article  
SET PrincipalAuthor='John Paul'  
WHERE Aid=3  
T1 commits

T2

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3

time

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3  
T2 commits

# Concurrency in a DBMS

## *Isolation Levels*

### REPEATABLE READ

- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object
- exclusive locks are released at the end of the transaction
- shared locks are released at the end of the transaction
- a transaction cannot change data that has been read by the current transaction, prior to its completion
- transactions exposed to the phantom problem

# Concurrency in a DBMS

## ~~Non-repeatable reads~~

- T2 reads the values: *Popescu Dan* and *John Paul* – under REPEATABLE READ

(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

UPDATE Article  
SET PrincipalAuthor='John Paul'  
WHERE Aid=3  
T1 commits

T2

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3

time

SELECT PrincipalAuthor  
FROM Article  
WHERE Aid=3  
T2 commits

# Concurrency in a DBMS

(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

## Phantoms

- T2 under REPEATABLE READ; T2 reads a set of tuples based on a search condition; T1 can add a row that alters the set of tuples, while T2 is still in progress
- T2 reads articles with Aids: (3, 4), (2, 3, 4) – under REPEATABLE READ

T1

INSERT INTO Article(Aid, ...)  
VALUES (2, ...)

T1 commits

T2

SELECT \*  
FROM Article  
WHERE Aid BETWEEN 1 AND 5

SELECT \*  
FROM Article  
WHERE Aid BETWEEN 1 AND 5  
T2 commits

time

# Concurrency in a DBMS

## ***Isolation Levels***

### SERIALIZABLE

- strongest isolation level
  - a transaction must acquire locks on objects before reading / writing them
  - a transaction also acquires locks on sets of objects that must remain unmodified
    - if a transaction T reads a set of objects based on a search predicate, this set cannot be changed by other transactions while T is in progress (if query *SELECT \* FROM Students WHERE Grade >= 8* is executed twice within a transaction, it must return the same answer set)
  - locks are held until the end of the transaction
- phantoms not possible

# Concurrency in a DBMS

## ~~Phantoms~~

- T2 reads articles with Aids: (3, 4), (3, 4) – under SERIALIZABLE

(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

INSERT INTO Article(Aid, ...)  
VALUES (2, ...)

T1 commits

T2

SELECT \*  
FROM Article  
WHERE Aid BETWEEN 1 AND 5

SELECT \*  
FROM Article  
WHERE Aid BETWEEN 1 AND 5  
T2 commits

time

# Concurrency in a DBMS

## Deadlocks

- e.g. T1, T2 – under REPEATABLE READ → deadlock (cycle of transactions waiting for one another to release a locked resource)  
(3, 'Artificial Intelligence for DNA', 'Popescu Dan', 2020, 5)  
(4, 'General AI characteristics', 'Kelly James', 2010, 4)

T1

@t = Types for article with Aid=3

UPDATE Article  
SET Type=@t +5  
WHERE Aid=3  
....

T1 commits

T2

@t = Types for article with Aid=3

UPDATE Article  
SET Type=@t +2  
WHERE Aid=3

T2 commits



# Concurrency in a DBMS

Concurrency Problems			
Isolation Levels	Dirty reads	Non-repeatable reads	Phantoms
Read Uncommittted	Yes	Yes	Yes
Read Committed	No	Yes	Yes
Repeatable Read	No	No	Yes
Serializable	No	No	No



# References:

- C.J. Date, *An Introduction to Database Systems (8th Edition)*, Addison-Wesley, 2003.
- H. Garcia-Molina, J. Ullman, J. Widom, *Database Systems: The Complete Book*, Prentice Hall Press, 2008.
- G. Hansen, J. Hansen, *Database Management And Design (2nd Edition)*, Prentice Hall, 1996.
- R. Ramakrishnan, J. Gehrke, *Database Management Systems*, McGraw- Hill, 2007.  
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- R. Ramakrishnan, J. Gehrke, *Database Management Systems (2nd Edition)*, McGraw-Hill, 2000.
- A. Silberschatz, H. Korth, S. Sudarshan, *Database System Concepts*, McGraw-Hill, 2010.  
<http://codex.cs.yale.edu/avi/db-book/>
- L. Țâmbulea, *Curs Baze de date*, Facultatea de Matematică și Informatică, UBB, 2013-2014.
- J. Ullman, J. Widom, *A First Course in Database Systems*,  
<http://infolab.stanford.edu/~ullman/fcdb.html>